

DataLens FAQs

Document ID: FAQ-PRD-DL-001 **Last Updated:** March 2024 **Owner:** DataLens Product Team **Category:** Product - DataLens

Getting Started

Q: What is DataLens?

A: DataLens is NovaTech's unified observability platform for monitoring your infrastructure, applications, and business metrics. Key features include: - Metrics collection and visualization - Log aggregation and search - Distributed tracing - Custom dashboards - Intelligent alerting - Natural language queries (beta)

Q: How do I send data to DataLens?

A:

Metrics (OpenTelemetry):

```
# otel-config.yaml
exporters:
  datalens:
    endpoint: https://ingest.datalens.novatech.com
    api_key: ${DATALENS_API_KEY}
```

Logs (Fluent Bit):

```
[OUTPUT]
Name http
Match *
Host ingest.datalens.novatech.com
Port 443
URI /api/v1/logs
Header Authorization Bearer ${DATALENS_API_KEY}
```

Application (SDK):

```
const datalens = require('@novatech/datalens');

datalens.init({
  apiKey: process.env.DATALENS_API_KEY,
  service: 'my-app'
});
```

Q: What data sources are supported?

A:

| Category | Sources |
|----------------|--|
| Infrastructure | AWS CloudWatch, Azure Monitor, GCP Stackdriver, Prometheus |
| Applications | OpenTelemetry, StatsD, custom SDKs |
| Logs | Fluent Bit, Fluentd, Logstash, syslog |
| Traces | OpenTelemetry, Jaeger, Zipkin |
| Databases | PostgreSQL, MySQL, MongoDB (query metrics) |
| Cloud | AWS, Azure, GCP native integrations |

Metrics

Q: How do I create a custom metric?

A:

Using SDK:

```
datalens.metrics.gauge('api.response_time', 150, {
  endpoint: '/users',
  method: 'GET'
});

datalens.metrics.counter('api.requests', 1, {
  status: '200'
});
```

Using StatsD:

```
echo "api.request.count:1|c" | nc -u localhost 8125
```

Q: What's the data retention for metrics?

A:

| Plan | Full Resolution | Downsampled |
|--------------|-----------------|-------------|
| Starter | 7 days | 30 days |
| Professional | 30 days | 1 year |
| Enterprise | 90 days | 3 years |

Downsampled data uses 1-hour resolution.

Q: How do I query metrics?

A: Use our query language (DQL):

```
-- Average CPU by host
SELECT avg(cpu.usage)
FROM metrics
WHERE service = 'api'
GROUP BY host
SINCE 1h

-- 95th percentile latency
SELECT percentile(http.latency, 95)
FROM metrics
WHERE environment = 'production'
SINCE 24h
INTERVAL 5m
```

Logs

Q: How do I search logs?

A: Use the log search bar or query:

```
level:error service:api message:"timeout"
```

Search operators: - field:value - Exact match - field:**pattern** - Wildcard - field:[1 TO 100] - Range - AND, OR, NOT - Boolean

Q: Can I set up log-based alerts?

A: Yes:

```
alert:
  name: error-rate-high
  type: log
  query: level:error service:api
  condition:
    threshold: 100
    window: 5m
  actions:
    - type: slack
      channel: "#alerts"
```

Q: How long are logs retained?

A:

| Plan | Hot Storage | Cold Storage |
|--------------|-------------|--------------|
| Starter | 3 days | 14 days |
| Professional | 15 days | 90 days |
| Enterprise | 30 days | 1 year+ |

Hot storage = fast search. Cold storage = archived, slower retrieval.

Tracing

Q: How do I enable distributed tracing?

A: Add the OpenTelemetry SDK:

```
const { NodeTracerProvider } = require('@opentelemetry/node');
const { DataLensExporter } = require('@novatech/datalens-otel');

const provider = new NodeTracerProvider();
provider.addSpanProcessor(
  new BatchSpanProcessor(new DataLensExporter({
    apiKey: process.env.DATALENS_API_KEY
  }))
);
provider.register();
```

Q: How do I trace across services?

A: Context propagates automatically with supported frameworks. For custom propagation:

```
const { propagation, context } = require('@opentelemetry/api');

// Inject into outgoing request
propagation.inject(context.active(), headers);

// Extract from incoming request
const ctx = propagation.extract(context.active(), headers);
```

Q: What's the trace sampling rate?

A: Default: 10% (configurable)

```
tracing:
  sampling:
    rate: 0.1 # 10%
    # Or use adaptive sampling
    adaptive:
      target_rate: 100 # traces per second
```

Dashboards

Q: How do I create a dashboard?

A:

1. Go to **DataLens** → **Dashboards** → **New**
 2. Add widgets (graphs, stats, tables)
 3. Configure queries for each widget
 4. Arrange layout
 5. Save and share
-

Q: Can I share dashboards?

A: Yes:

- **Internal sharing:** Add team members as viewers/editors
 - **Public link:** Create read-only shareable link
 - **Embed:** Generate iframe embed code
 - **Export:** PDF or PNG snapshot
-

Q: Are there pre-built dashboards?

A: Yes! Templates available:

- Kubernetes Overview
- AWS Infrastructure
- Application Performance
- Database Monitoring
- Business Metrics

Go to **Dashboards** → **Templates** to import.

Alerting

Q: How do I create an alert?

A:

UI: 1. Go to **Alerts** → **New Alert** 2. Define query and threshold 3. Set notification channels 4. Configure schedule

YAML:

```
alert:
  name: high-latency
  query: |
    SELECT avg(http.latency) FROM metrics
    WHERE service = 'api'
  condition:
    operator: ">"
    threshold: 500
    duration: 5m
  notifications:
    - type: pagerduty
      severity: high
```

Q: What notification channels are supported?

A:

| Channel | Setup Required |
|-----------------|----------------------|
| Email | Built-in |
| Slack | Workspace connection |
| PagerDuty | Integration key |
| OpsGenie | Integration key |
| Webhook | URL configuration |
| Microsoft Teams | Connector URL |

Q: How do I avoid alert fatigue?

A: Best practices:

1. **Use appropriate thresholds** - Alert on business impact
 2. **Add duration requirements** - Avoid transient spikes
 3. **Group related alerts** - Reduce noise
 4. **Use severity levels** - Page only for critical
 5. **Enable auto-resolve** - Clear when condition ends
-

Natural Language Queries (Beta)

Q: How do I use natural language queries?

A: Just type your question:

"What caused errors to spike yesterday at 3pm?"
"Show me the slowest endpoints this week"
"Compare today's traffic to last Tuesday"

Q: How accurate is NLQ?

A: Current accuracy: ~85% for common queries. The system:
- Shows confidence level
- Displays interpreted query
- Allows query refinement

If results seem wrong, click "Show query" to verify interpretation.

Pricing

Q: How is DataLens billed?

A:

| Component | Unit |
|-----------|-------------------------------|
| Metrics | Per million data points/month |
| Logs | Per GB ingested |
| Traces | Per million spans |
| Retention | Additional for extended |

Check usage at **Settings → Usage & Billing**.

Q: What's included in each plan?

A:

| Feature | Starter | Professional | Enterprise |
|------------|----------|--------------|------------|
| Metrics | 10M/mo | 100M/mo | Unlimited |
| Logs | 10 GB/mo | 100 GB/mo | Unlimited |
| Traces | 1M/mo | 10M/mo | Unlimited |
| Users | 5 | 25 | Unlimited |
| Dashboards | 10 | Unlimited | Unlimited |
| Alerts | 10 | 100 | Unlimited |
| NLQ | No | Yes | Yes |

Troubleshooting

Q: Data not appearing in DataLens?

A: Check: 1. API key is valid 2. Data is being sent (check logs) 3. Time range is correct 4. Data retention hasn't expired

Debug:

```
curl -X POST https://ingest.datalens.novatech.com/api/v1/test \  
-H "Authorization: Bearer $DATALENS_API_KEY"
```

Q: Dashboard is slow?

A:

1. Reduce time range
 2. Add filters to queries
 3. Use appropriate aggregations
 4. Reduce widget count
 5. Enable caching
-

Q: Alerts not firing?

A: Verify: 1. Alert is enabled 2. Query returns data 3. Threshold is appropriate
4. Notification channel is configured

Test with: **Alerts → [Your Alert] → Test**

Integration

Q: Does DataLens integrate with CloudForge?

A: Yes! Native integration:
- Automatic metric collection
- Environment health dashboards
- Deployment annotations
- Cost metrics

Enable in CloudForge: **Settings → Integrations → DataLens**

Q: Can I export data from DataLens?

A:

```
# Export metrics
datalens export metrics \
--query "SELECT * FROM metrics WHERE service='api'" \
--format csv \
--output metrics.csv

# Export to S3
datalens export metrics \
--query "SELECT * FROM metrics" \
--destination s3://my-bucket/exports/
```

Contact

- **Documentation:** docs.novatech.com/datalens
 - **Support:** support@novatech.com
 - **Slack:** #datalens-help
-

Last reviewed: March 2024