

# CloudForge Architecture Overview

**Document ID:** PRD-CF-002 **Last Updated:** 2024-02-15 **Owner:** CloudForge Engineering **Classification:** Internal

---

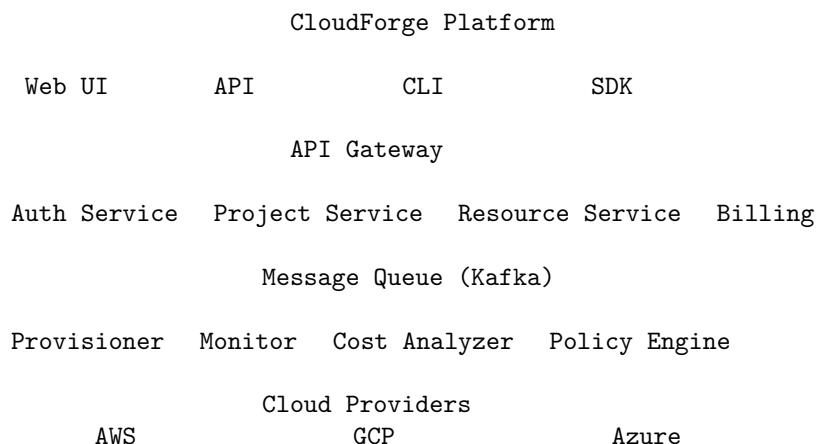
## Introduction

CloudForge is NovaTech's flagship cloud infrastructure management platform. This document provides an architectural overview for internal teams, partners, and customers seeking to understand the system's design.

---

## System Architecture

### High-Level Components



---

## Core Services

### API Gateway

- Entry point for all client requests

- Rate limiting and throttling
- Request routing and load balancing
- API versioning support (v1, v2)
- Technology: Kong + custom plugins

### **Authentication Service**

- User authentication via Okta integration
- API key management
- Service-to-service authentication (mTLS)
- RBAC policy enforcement
- Technology: Go, PostgreSQL

### **Project Service**

- Multi-tenant project management
- Organization hierarchy
- Team membership and permissions
- Audit logging
- Technology: Go, PostgreSQL

### **Resource Service**

- Cloud resource lifecycle management
- State tracking and reconciliation
- Dependency resolution
- Drift detection
- Technology: Go, PostgreSQL, Redis

### **Provisioner Service**

- Infrastructure provisioning engine
- Terraform execution runtime
- Multi-cloud provider adapters
- Parallel resource creation
- Technology: Go, Terraform

### **Monitor Service**

- Real-time resource monitoring
- Health checks and alerting
- Metrics collection and aggregation

- Integration with external monitoring
- Technology: Go, InfluxDB, Prometheus

### Cost Analyzer Service

- Cloud cost aggregation
- Budget tracking and alerts
- Cost allocation by project/team
- Optimization recommendations
- Technology: Python, PostgreSQL, BigQuery

### Policy Engine

- Infrastructure policy enforcement
  - Compliance rule evaluation
  - Pre-deployment validation
  - Real-time policy monitoring
  - Technology: Go, Open Policy Agent
- 

## Data Architecture

### Primary Database (PostgreSQL)

- User and organization data
- Project configurations
- Resource state and metadata
- Audit logs

**Cluster Configuration:** - Primary + 2 read replicas - Automated failover - Point-in-time recovery (30 days) - Cross-region replication

### Cache Layer (Redis)

- Session data
- API response caching
- Rate limiting counters
- Resource state cache

**Cluster Configuration:** - 6-node cluster with sentinel - Automatic failover - Cache invalidation on writes

## Time-Series Database (InfluxDB)

- Resource metrics
- Performance data
- Cost time-series data

## Message Queue (Kafka)

- Async task processing
- Event streaming
- Service communication
- Audit event stream

**Topics:** - `resource.events` - Resource lifecycle events - `provisioner.tasks` - Provisioning jobs - `monitor.metrics` - Monitoring data - `audit.logs` - Audit trail

---

## Infrastructure

### Cloud Provider: AWS (Primary)

- Region: us-west-2 (primary), us-east-1 (DR)
- EKS for container orchestration
- RDS for PostgreSQL
- ElastiCache for Redis
- MSK for Kafka

## Kubernetes Architecture

- Production cluster: 50+ nodes
- Node pools by workload type
- Horizontal Pod Autoscaler
- Pod Disruption Budgets

## CDN and Edge

- CloudFront for static assets
- Global edge for API acceleration
- DDoS protection via AWS Shield

## Security Architecture

### Network Security

- VPC isolation per environment
- Private subnets for all services
- WAF for API protection
- Network policies (Calico)

### Data Security

- Encryption at rest (AES-256)
- Encryption in transit (TLS 1.3)
- Customer-managed keys (BYOK) option
- Data residency controls

### Access Control

- RBAC with custom roles
  - Organization-level policies
  - Project-level permissions
  - Audit logging for all access
- 

## Scalability

### Auto-Scaling Policies

Service	Min	Max	Scale Metric
API Gateway	5	50	Request rate
Resource Service	3	30	CPU
Provisioner	10	100	Queue depth
Monitor	5	25	Memory

### Performance Targets

Metric	Target	Current
API Latency (p99)	<200ms	150ms

---

Metric	Target	Current
Provisioning Time	<5min	3.2min
UI Load Time	<2s	1.5s
Availability	99.9%	99.95%

---

## Disaster Recovery

### RTO/RPO

- RTO (Recovery Time Objective): 4 hours
- RPO (Recovery Point Objective): 1 hour

### Backup Strategy

- Database: Continuous replication + daily snapshots
- Configuration: Version controlled (Git)
- Secrets: Vault with HA backend

### Failover Process

1. Automated health checks detect failure
  2. DNS failover to DR region (Route 53)
  3. DR cluster promoted to primary
  4. Data sync verification
  5. Traffic shift complete
- 

## Integration Points

### Inbound Integrations

System	Protocol	Purpose
Customer Apps	REST API	Resource management
CI/CD Systems	REST API / Webhooks	Automated deployments
Monitoring Tools	Prometheus	Metrics export
SIEM	Syslog / API	Security logs

## Outbound Integrations

System	Protocol	Purpose
AWS	AWS SDK	Cloud provisioning
GCP	Google Cloud SDK	Cloud provisioning
Azure	Azure SDK	Cloud provisioning
Slack	Webhooks	Notifications
PagerDuty	API	Incident alerts

---

## Future Architecture

### Planned Improvements

- Multi-region active-active (Q3 2024)
  - GraphQL API support (Q4 2024)
  - Edge computing integration (2025)
  - AI-powered optimization engine (2025)
- 

*Related Documents:* *Getting Started Guide (PRD-CF-001), API Reference (PRD-CF-010), Security Whitepaper (PRD-CF-020)*