

CloudForge Monitoring Integration Guide

Document ID: PRD-CF-045 **Last Updated:** 2024-03-01 **Owner:** CloudForge Product Team **Classification:** Public

Overview

CloudForge integrates seamlessly with DataLens and third-party monitoring tools to provide comprehensive observability for your infrastructure. This guide covers built-in monitoring, integrations, and best practices.

Built-in Monitoring

Default Metrics

CloudForge automatically collects metrics for all resources:

Compute Metrics

Metric	Description	Interval
cpu_utilization	CPU usage percentage	1 min
memory_utilization	Memory usage percentage	1 min
disk_read_bytes	Disk read throughput	1 min
disk_write_bytes	Disk write throughput	1 min
network_in_bytes	Network ingress	1 min
network_out_bytes	Network egress	1 min

Service Metrics

Metric	Description	Interval
request_count	HTTP requests	1 min
request_latency	Response time	1 min
error_rate	Error percentage	1 min
active_connections	Current connections	1 min
replica_count	Running replicas	1 min

Database Metrics

Metric	Description	Interval
connections	Active connections	1 min
cpu_utilization	Database CPU	1 min
storage_used	Storage consumption	5 min
read_iops	Read operations	1 min
write_iops	Write operations	1 min
replication_lag	Replica lag seconds	1 min

DataLens Integration

Enable Integration

```
# environment.yaml
monitoring:
  datalens:
    enabled: true
    workspace: production

  metrics:
    enabled: true
    retention: 15d

  logs:
    enabled: true
    retention: 7d

  traces:
    enabled: true
    sampling_rate: 0.1
```

Automatic Dashboard

When integration is enabled, CloudForge creates dashboards automatically:

- **Environment Overview** - Health status, resource usage
- **Service Performance** - Request rates, latency, errors
- **Database Monitoring** - Connections, queries, performance
- **Cost Analysis** - Spend breakdown, trends

Custom Metrics

Send custom metrics from your applications:

```
const datalens = require('@novatech/datalens');

// Initialize with CloudForge context
datalens.init({
  source: 'cloudforge',
  environment: process.env.CLOUDFORGE_ENV
});

// Send custom metric
datalens.gauge('order.processing_time', 1250, {
  order_type: 'subscription',
  region: 'us-west'
});
```

Alerting

Built-in Alerts

CloudForge provides default alert rules:

Alert	Condition	Default Threshold
High CPU	CPU > threshold for 5m	80%
High Memory	Memory > threshold for 5m	85%
High Error Rate	Errors > threshold for 2m	5%
Service Down	Health check failing	3 consecutive
Database Connection High	Connections > threshold	80% of max
Disk Space Low	Disk usage > threshold	85%

Custom Alerts

```
# alerts.yaml
alerts:
  - name: high-latency
    description: API latency too high
    metric: request_latency_p99
    condition:
      operator: ">"
```

```

        threshold: 500
        duration: 5m
severity: warning
notifications:
  - type: slack
    channel: "#alerts"
  - type: pagerduty
    severity: low

- name: error-spike
  description: Sudden increase in errors
  metric: error_rate
  condition:
    operator: ">"
    threshold: 10
    duration: 2m
  severity: critical
  notifications:
    - type: pagerduty
      severity: high
    - type: email
      recipients:
        - oncall@novatech.com

```

Alert Routing

```

notifications:
  routes:
    - match:
        severity: critical
        receivers: [pagerduty, slack-critical]

    - match:
        severity: warning
        receivers: [slack-warnings]

    - match:
        environment: production
        receivers: [prod-team]

  receivers:
    - name: pagerduty
      type: pagerduty
      routing_key: ${PAGERDUTY_KEY}

```

```
- name: slack-critical
  type: slack
  webhook: ${SLACK_CRITICAL_WEBHOOK}
  channel: "#critical-alerts"
```

Third-Party Integrations

Prometheus

Export metrics to your Prometheus instance:

```
monitoring:
  prometheus:
    enabled: true
    endpoint: https://prometheus.internal.novatech.com
    credentials:
      type: bearer
      token: ${PROMETHEUS_TOKEN}
    scrape_interval: 30s
```

Grafana

CloudForge provides Grafana dashboards:

```
monitoring:
  grafana:
    enabled: true
    url: https://grafana.internal.novatech.com
    api_key: ${GRAFANA_API_KEY}
    dashboards:
      - cloudforge-overview
      - service-performance
      - database-health
```

Datadog

Send metrics to Datadog:

```
monitoring:
  datadog:
```

```
    enabled: true
    api_key: ${DATADOG_API_KEY}
    site: datadoghq.com
    tags:
      - env:production
      - team:platform
```

New Relic

```
monitoring:
  newrelic:
    enabled: true
    license_key: ${NEWRELIC_LICENSE_KEY}
    account_id: ${NEWRELIC_ACCOUNT_ID}
```

PagerDuty

Direct integration for incident management:

```
monitoring:
  pagerduty:
    enabled: true
    routing_key: ${PAGERDUTY_ROUTING_KEY}
    severity_mapping:
      critical: critical
      warning: warning
      info: info
```

Logging

Log Configuration

```
logging:
  level: info
  format: json

  outputs:
    - type: datalens
      enabled: true

    - type: s3
```

```

    enabled: true
    bucket: novatech-logs
    prefix: cloudforge/

    - type: cloudwatch
      enabled: true
      log_group: /cloudforge/production

```

Structured Logging

CloudForge adds context to all logs:

```
{
  "timestamp": "2024-07-25T15:00:00Z",
  "level": "info",
  "message": "Request completed",
  "service": "api",
  "environment": "production",
  "trace_id": "abc123",
  "span_id": "xyz789",
  "request": {
    "method": "POST",
    "path": "/api/v1/users",
    "status": 201,
    "duration_ms": 145
  }
}
```

Log Queries

Search logs via CLI:

```

# Recent errors
cloudforge logs --env production --level error --since 1h

# Specific service
cloudforge logs --env production --service api --filter "timeout"

# Trace correlation
cloudforge logs --env production --trace-id abc123

```

Distributed Tracing

Enable Tracing

```
tracing:
  enabled: true
  provider: datalens
  sampling:
    rate: 0.1 # 10% of requests
    # Or adaptive sampling
    adaptive:
      target_rate: 100 # traces per second
```

Trace Context Propagation

CloudForge automatically propagates trace context:

```
tracing:
  propagation:
    - tracecontext # W3C standard
    - baggage
    - b3           # Zipkin format
```

Custom Spans

Add custom spans in your application:

```
const tracer = require('@novatech/tracing');

async function processOrder(order) {
  return tracer.startActiveSpan('process_order', async (span) => {
    span.setAttribute('order_id', order.id);

    // Processing logic
    await validateOrder(order);
    await chargePayment(order);
    await fulfillOrder(order);

    span.end();
  });
}
```

Health Checks

Service Health Checks

```
services:
  - name: api
    health_check:
      path: /health
      port: 8080
      interval: 30s
      timeout: 5s
      healthy_threshold: 2
      unhealthy_threshold: 3
```

Readiness vs Liveness

```
services:
  - name: api
    probes:
      liveness:
        path: /healthz
        initial_delay: 30s
        period: 10s
      readiness:
        path: /ready
        initial_delay: 5s
        period: 5s
```

Health Check Response

Standard health check response format:

```
{
  "status": "healthy",
  "version": "2.5.0",
  "uptime": 86400,
  "checks": {
    "database": {
      "status": "healthy",
      "latency_ms": 5
    },
    "cache": {
      "status": "healthy",
      "latency_ms": 1
    }
  }
}
```

```
        },
        "external_api": {
            "status": "degraded",
            "latency_ms": 250,
            "message": "Slow response"
        }
    }
}
```

Deployment Monitoring

Deployment Events

CloudForge creates annotations for deployments:

```
monitoring:
  annotations:
    deployments: true
    config_changes: true
    scaling_events: true
```

Deployment Metrics

Metric	Description
deployment_duration	Time to complete deployment
deployment_rollback_count	Rollbacks in period
deployment_success_rate	Successful deployments %

Change Correlation

Automatically correlate issues with recent changes:

```
Alert: Error rate spike at 15:05
Possible cause: Deployment at 15:02 (api v2.5.0 → v2.5.1)
```

Best Practices

Dashboard Design

1. Start with **RED method:**

- Rate (requests/second)
- Errors (error rate)
- Duration (latency)

2. Add **USE method for resources:**

- Utilization
- Saturation
- Errors

3. Include business metrics:

- Orders processed
- Revenue impact
- User actions

Alert Configuration

1. Alert on symptoms, not causes
2. Set meaningful thresholds
3. Include runbook links
4. Avoid alert fatigue
5. Review and tune regularly

Retention Strategy

Data Type	Hot Storage	Cold Storage
Metrics	15 days	1 year
Logs	7 days	90 days
Traces	7 days	30 days

Related Documents: DataLens Integration (PRD-DL-040), Alerting Guide (PRD-CF-046), Troubleshooting (PRD-CF-070)